



UNIVERSITÀ
DEGLI STUDI
DI PADOVA



DIPARTIMENTO
DI INGEGNERIA
DELL'INFORMAZIONE

Lecture 18

Bayesian games

Thomas Marchioro

December 6, 2023

- Static games **of complete information**:
 - Players move simultaneously/without knowing each other's move
 - Game ends after a single interaction
- Dynamic games **of complete information**:
 - Sequential games: perfect information, each player knows exactly where they are in the extensive form tree
 - Multistage/repeated games: sequences of games (static or dynamic)


- Remember: Complete information means
 - Every player knows who are the other players
 - Every player knows the possible moves available to other players
 - Every player has full knowledge of other players' utility function → We drop this assumption
 - today, we see the case where some players might just have an estimate of the “**type**” of the other player
 - Everybody knows that everybody knows all this information
- **Bayesian games**: we do not know exactly the utility function of other players, we consider different possibilities
- Drop the assumption of complete information

Example: Lazy student's problem

- Suppose you decided to make a project for your Game Theory course, and you need to decide whether to put some actual effort (E), or to make a halfhearted project (H)
- If you put effort, you are guaranteed to get a good score on the project, which gives you utility +1
- However, ideally you would like to slack off and still get a good grade, which gives you utility +2; on the other hand, there is the possibility that your project gets a bad grade (payoff -1)
- It is up to your professor whether to give a good grade (G) to your halfhearted project or to flunk you (F)

Example: Lazy student's problem


- Problem is: you know your professor's possible moves (F or G), but you do not know his utility function
- It is reasonable to assume that the professor is always happy to see a good project, so his payoff if you put effort is $+1$



- Consider two possible scenarios:
 - *Scenario 1*: your professor is a **lenient** professor and always prefers to avoid flunking students (grading is a very boring activity, he does not want to do that twice). Playing F gives him a payoff of -1 , whereas G gives him a payoff of 0
 - *Scenario 2*: your professor is a **strict** professor who believes that students who are slacking off deserve to get flunked. Playing F gives him a payoff of $+1$, and G still gives him a payoff of 0

Example: Lazy student's problem


- *Scenario 1: lenient professor case*



- SPE (H, G) found via backward induction
- (E, F) is a NE but not a SPE: non-credible threat

Example: Lazy student's problem

■ Scenario 2: **strict** professor case



- SPE (E, F) found via backward induction (also, only NE of the game)


Example: Lazy student's problem

- Same problem as before: you do not know whether your professor is lenient or strict
- Does that mean that you cannot make a rational decision?
- You may think of using a conservative approach and always assume the worst-case scenario (in this case, assuming the strict-professor scenario)
- However, that is not how people make decisions: how they behave depends on their estimate of each scenario
 - If you are 99% confident that your professor is lenient, you will be more inclined to make a halfhearted project
 - If you are 99% confident that your professor is strict, you are more likely to put effort
 - Can we generalize this for any estimate p ?

- Bayesian games are games of **incomplete information**, where some players do not have complete knowledge on the other players' utility
- This incomplete information, however, still needs to be modeled somehow
- Bayesian game assumptions:
 - Each player i has a set of possible types
 $T_i = \{\text{type 1, type 2, } \dots\}$ (e.g., $T_P = \{\text{lenient, strict}\}$)
 - This set can be a singleton for some players (if all players can have only one type \rightarrow back to complete information)
 - The type of each player t_1, \dots, t_n is determined by Nature's move *at the beginning of the game*
 - Each player knows only his/her own type
 - Other player's type is unknown


- Bayesian games can be seen as a weird kind of dynamic games that is played as follows:
 - Nature draws a type vector (t_1, \dots, t_n) among all the possible combinations of players' types
 - Nature reveals type t_i only to player i (this is captured by information sets)
 - Players choose their actions
 - Final payoffs are computed
- This is a dynamic game where the players do not know Nature's move in its entirety

- Nature chooses at the beginning whether P is being “lenient” or “strict”
 - S does not know about Nature’s choice: only one information set
 - P knows about Nature’s choice: one information set for “lenient” and one for “strict”



- Players know their types but not their opponents'
- How can they find best responses?
 - They create beliefs about these types
 - *Assumption*: players do not precisely know the types of their opponents but they have an estimate of those
 - In other words, they know the **probability distribution** of the opponents' types (this is common knowledge!)
 - This means that all players know that player i is of type t_1 with probability $p_i(t_1)$, t_2 with probability $p_i(t_2)$, and so on
 - Player i knows about his/her true type but also knows the opponents' estimates $p_i(t_1), p_i(t_2), \dots$
 - This is called the **common prior assumption**


- Suppose your estimate is that the professor is lenient with probability p and strict with probability $1 - p$
- The **extensive form** of the Bayesian game is as follows



- The lazy student's problem is a more general game called **"Bayesian entry game"**
- Original version: incumbent (**player 1**) vs. outsider (player 2)
- The outsider decides whether to enter the market (E) or stay out (O)
- The incumbent decides whether to accept the outsider (A) or fight (F)
- The incumbent could be of two types: **"reasonable"** (does not like to fight) or **"crazy"** (likes to fight)
- The incumbent knows its type, the outsider estimates reasonable/crazy with probabilities (**p** , **$1 - p$**)

Bayesian entry game


- **Extensive form** of the Bayesian entry game:



- The normal form of Bayesian games can be inferred from the extensive form as we did in dynamic games
- Let us start by listing strategies: this can be done by considering all combinations possible moves in each information set
- In the Bayesian Entry Game:
 - Player 1 only has one information set (does not know Nature's move, which decide player 2's type): only one move needed to define pure strategies
 - Player 2 has two information sets (knows his/her type): his/her pure strategies are all possible combination of moves in these 2 information sets

- In Bayesian games, the concept of strategy is further extended: a strategy specifies what the player does for each type
 - Think of it as a “Dr. Jekyll and Mr. Hyde” situation
 - E.g., strategy **AF** for player 2 means “If I am a “reasonable” type, I accept the outsider; if I am a “crazy” type, I fight them”
 - This might seem strange, since players know their actual type; however, they need to do this reasoning in order to predict the opponents’ strategy
- *Example:* In the Bayesian Entry Game:
 - Player 1 has two strategies: $\{O, E\}$
 - Player 2 has four strategies: $\{AA, AF, FA, FF\}$
- Each pair of pure strategies determines how the game is played, which also depends on Nature’s choice

Bayesian entry game



- For example, if the game is played as (E, AF)
 - **player 1** gets: $p \cdot 1 + (1 - p) \cdot (-1) = 2p - 1$
 - **player 2** gets: $p \cdot 1 + (1 - p) \cdot 2 = 2 - p$

- We can determine player 2's best responses against O (anything)
- We can determine player 1's best response against AA and FF (when 2 always plays the same regardless of its type)
- So we know that (O,FF) is a NE regardless of p
- However, we need to know p in order to find *all* NE

		Player 2			
		AA	AF	FA	FF
Player 1	O	0, 2	0, 2	0, 2	0, 2
	E	1, 1	$2p-1, 2-p$	$1-2p, 1-2p$	$-1, 2-3p$

- For $p = 0$ (The outsider is sure that the incumbent is **crazy**)

		Player 2			
		AA	AF	FA	FF
Player 1	O	0, 2	0, 2	0, 2	0, 2
	E	1, 1	-1, 2	1, 1	-1, 2

- (O, AF) and (O, FF) are NE
- Payoffs of player 2 only reflect the choice of the crazy type (choices of reasonable type are neglected)
- **Interpretation:** the outsider (player 1) is 100% sure that the incumbent is crazy, so the moves of a “reasonable” player 2 are ignored
- Therefore, the NE are all joint strategies (O, *F)

Bayesian entry game, NE

- For $p = 1$ (The outsider is sure that the incumbent is reasonable)

		Player 2			
		AA	AF	FA	FF
Player 1	O	0, 2	0, 2	0, 2	0, 2
	E	1, 1	1, 1	-1, -1	-1, -1

- (E, AA), (E, FA) and (E, AF) are Nash equilibria
- **Interpretation:** the outsider is 100% sure that the incumbent is reasonable, so the crazy-incumbent moves are ignored
- NE are (E, A*) and (O, F*)

Bayesian entry game, NE

- For $p = 2/3$

		Player 2			
		AA	AF	FA	FF
Player 1	O	0, 2	0, 2	0, 2	0, 2
	E	1, 1	1/3, 4/3	-1/3, -1/3	-1, 0

- NE: (E, AF), (O, FA), (O, FF)
- **Interpretation:** The outsider thinks that the incumbent is likely to be **reasonable**, but there is still some probability of them being **crazy**. Therefore, joint strategies like (E, AA) where also the crazy type is accepting the outsider are not NE.

Bayesian games: definition

- Up to now, our normal-form representation of a game included:
 - Set of players $\{1, \dots, n\}$
 - Strategy sets S_1, \dots, S_n
 - Utility functions $u_i : S_1 \times \dots \times S_n \rightarrow \mathbb{R}$ (for $i = 1, \dots, n$)
- In Bayesian games we need three more ingredients:
 - 1 Type space** of each player T_i (for $i = 1, \dots, n$)
 - 2 Type-dependent utilities:** define $u_i(a_1, \dots, a_n, t_i)$ for each $t_i \in T_i$
 - 3 Beliefs about players' types:** a probability distribution ϕ_i defined over types for each player

- Consider a static Bayesian game: n players, each player's (pure) strategy is just an action
 - 1 Player i 's type is $t_i \in T_i$, chosen by Nature for each player from 1 to n through the joint **prior probability distribution** $\phi(t_1, \dots, t_n)$, where $\phi : T_1 \times \dots \times T_n \rightarrow [0, 1]$
 - This prior is **common knowledge** among players
 - 2 Player i knows his **private values** of his/her utility function $u_i(a_1, \dots, a_n, t_i)$
 - Other players know **common values** of i 's utility function

- For example, suppose player i can have two different payoff functions $u_{i,A}(a_i, a_{-i})$ and $u_{i,B}(a_i, a_{-i})$
 - We represent this by setting a type space $T_i = \{t_A, t_B\}$ and imposing

$$u_{i,j}(a_i, a_{-i}) = u_i(a_i, a_{-i}, t_j)$$

- Types can be used to limit available actions
 - If a player has actions $\{F, G, H\}$, but H is permitted only with probability q , we define types t_A and t_B
 - t_A and t_B have respective probabilities q and $1 - q$
 - In both cases $\{F, G, H\}$ are feasible actions, but all payoffs of move H under type t_B are $-\infty$

- 3 Types can be correlated; they are independent if

$$\phi(t_1, \dots, t_n) = \phi(t_1) \cdots \phi(t_n)$$

- i knows his/her own type, but not others' (t_{-i})
 - He/she estimates that as $\phi_i(t_{-i}|t_i)$, exploiting the correlation (**belief** is derived from conditional probability)
 - Our assumption of the prior being common knowledge equals to **perfect information**
 - In the case of incomplete and imperfect info., belief $\phi_i(t_{-i}|t_i)$ may even be wrong and have nothing to do with the true prior

Questions?